

Predicting Default Payments of Credit Card Clients

Sarbajeet Biswal

7 June 2018

```
library(dplyr)
```

```
## Warning: package 'dplyr' was built under R version 3.4.3
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.4.3
```

```
library(gridExtra)
```

```
## Warning: package 'gridExtra' was built under R version 3.4.2
```

```
##  
## Attaching package: 'gridExtra'
```

```
## The following object is masked from 'package:dplyr':  
##  
##   combine
```

```
library(arules)
```

```
## Warning: package 'arules' was built under R version 3.4.4
```

```
## Loading required package: Matrix
```

```
##  
## Attaching package: 'arules'
```

```
## The following object is masked from 'package:dplyr':  
##  
##      recode
```

```
## The following objects are masked from 'package:base':  
##  
##      abbreviate, write
```

```
df = read.csv("C:/Users/Administrator/Desktop/PYTHON/default_of_credit_card_clients.csv")  
  
df1 = df  
names(df1)
```

```
## [1] "ID" "LIMIT_BAL"  
## [3] "SEX" "EDUCATION"  
## [5] "MARRIAGE" "AGE"  
## [7] "PAY_0" "PAY_2"  
## [9] "PAY_3" "PAY_4"  
## [11] "PAY_5" "PAY_6"  
## [13] "BILL_AMT1" "BILL_AMT2"  
## [15] "BILL_AMT3" "BILL_AMT4"  
## [17] "BILL_AMT5" "BILL_AMT6"  
## [19] "PAY_AMT1" "PAY_AMT2"  
## [21] "PAY_AMT3" "PAY_AMT4"  
## [23] "PAY_AMT5" "PAY_AMT6"  
## [25] "default.payment.next.month"
```

Data cleaning


Checking null values

```
colSums(is.na(df1))
```

```
##          ID          LIMIT_BAL
##          0          0
##          SEX        EDUCATION
##          0          0
##          MARRIAGE    AGE
##          0          0
##          PAY_0       PAY_2
##          0          0
##          PAY_3       PAY_4
##          0          0
##          PAY_5       PAY_6
##          0          0
##          BILL_AMT1    BILL_AMT2
##          0          0
##          BILL_AMT3    BILL_AMT4
##          0          0
##          BILL_AMT5    BILL_AMT6
##          0          0
##          PAY_AMT1     PAY_AMT2
##          0          0
##          PAY_AMT3     PAY_AMT4
##          0          0
##          PAY_AMT5     PAY_AMT6
##          0          0
## default.payment.next.month
##          0
```

Checking outliers

```
boxplot(df1)
```



```
summary(df1)
```

```

##          ID          LIMIT_BAL          SEX          EDUCATION
##  Min.      :    1    Min.      : 10000    Min.      :1.000    Min.      :0.000
##  1st Qu.: 7501    1st Qu.:  5000    1st Qu.:1.000    1st Qu.:1.000
##  Median :15000    Median : 14000    Median :2.000    Median :2.000
##  Mean   :15000    Mean   : 167484    Mean   :1.604    Mean   :1.853
##  3rd Qu.:22500    3rd Qu.: 240000    3rd Qu.:2.000    3rd Qu.:2.000
##  Max.    :30000    Max.    :1000000    Max.    :2.000    Max.    :6.000
##  MARRIAGE          AGE          PAY_0          PAY_2
##  Min.      :0.000    Min.      :21.00    Min.      : -2.0000    Min.      : -2.0000
##  1st Qu.:1.000    1st Qu.:28.00    1st Qu.: -1.0000    1st Qu.: -1.0000
##  Median :2.000    Median :34.00    Median : 0.0000    Median : 0.0000
##  Mean   :1.552    Mean   :35.49    Mean   : -0.0167    Mean   : -0.1338
##  3rd Qu.:2.000    3rd Qu.:41.00    3rd Qu.: 0.0000    3rd Qu.: 0.0000
##  Max.    :3.000    Max.    :79.00    Max.    : 8.0000    Max.    : 8.0000
##  PAY_3          PAY_4          PAY_5          PAY_6
##  Min.      : -2.0000    Min.      : -2.0000    Min.      : -2.0000    Min.      : -2.0000
##  1st Qu.: -1.0000    1st Qu.: -1.0000    1st Qu.: -1.0000    1st Qu.: -1.0000
##  Median : 0.0000    Median : 0.0000    Median : 0.0000    Median : 0.0000
##  Mean   : -0.1662    Mean   : -0.2207    Mean   : -0.2662    Mean   : -0.2911
##  3rd Qu.: 0.0000    3rd Qu.: 0.0000    3rd Qu.: 0.0000    3rd Qu.: 0.0000
##  Max.    : 8.0000    Max.    : 8.0000    Max.    : 8.0000    Max.    : 8.0000
##  BILL_AMT1          BILL_AMT2          BILL_AMT3          BILL_AMT4
##  Min.      : -165580    Min.      : -69777    Min.      : -157264    Min.      : -170000
##  1st Qu.:   3559    1st Qu.:   2985    1st Qu.:   2666    1st Qu.:   2327
##  Median :  22382    Median :  21200    Median :   20089    Median :   19052
##  Mean   :  51223    Mean   :  49179    Mean   :   47013    Mean   :   43263
##  3rd Qu.:  67091    3rd Qu.:  64006    3rd Qu.:   60165    3rd Qu.:   54506
##  Max.    : 964511    Max.    :983931    Max.    :1664089    Max.    : 891586
##  BILL_AMT5          BILL_AMT6          PAY_AMT1          PAY_AMT2
##  Min.      : -81334    Min.      : -339603    Min.      :      0    Min.      :      0
##  1st Qu.:   1763    1st Qu.:   1256    1st Qu.:   1000    1st Qu.:    833
##  Median :  18105    Median :   17071    Median :   2100    Median :   2009
##  Mean   :  40311    Mean   :   38872    Mean   :   5664    Mean   :   5921
##  3rd Qu.:  50191    3rd Qu.:   49198    3rd Qu.:   5006    3rd Qu.:   5000
##  Max.    :927171    Max.    : 961664    Max.    :873552    Max.    :1684259
##  PAY_AMT3          PAY_AMT4          PAY_AMT5          PAY_AMT6
##  Min.      :      0    Min.      :      0    Min.      :    0.0    Min.      :    0.0
##  1st Qu.:   390    1st Qu.:   296    1st Qu.:   252.5    1st Qu.:   117.8
##  Median :  1800    Median :   1500    Median :  1500.0    Median :  1500.0
##  Mean   :  5226    Mean   :   4826    Mean   :  4799.4    Mean   :  5215.5
##  3rd Qu.:  4505    3rd Qu.:   4013    3rd Qu.:  4031.5    3rd Qu.:  4000.0
##  Max.    :896040    Max.    :621000    Max.    :426529.0    Max.    :528666.0
##  default.payment.next.month
##  Min.      :0.0000
##  1st Qu.:0.0000
##  Median :0.0000
##  Mean   :0.2212
##  3rd Qu.:0.0000
##  Max.    :1.0000

```

Structure of data

```
str(df1)
```

```
## 'data.frame':    30000 obs. of  25 variables:
## $ ID                : int  1 2 3 4 5 6 7 8 9 10 ...
## $ LIMIT_BAL         : int  20000 120000 90000 50000 50000 50000 500000 100000 140
000 20000 ...
## $ SEX               : int  2 2 2 2 1 1 1 2 2 1 ...
## $ EDUCATION         : int  2 2 2 2 2 1 1 2 3 3 ...
## $ MARRIAGE          : int  1 2 2 1 1 2 2 2 1 2 ...
## $ AGE               : int  24 26 34 37 57 37 29 23 28 35 ...
## $ PAY_0             : int  2 -1 0 0 -1 0 0 0 0 -2 ...
## $ PAY_2             : int  2 2 0 0 0 0 0 -1 0 -2 ...
## $ PAY_3             : int  -1 0 0 0 -1 0 0 -1 2 -2 ...
## $ PAY_4             : int  -1 0 0 0 0 0 0 0 0 -2 ...
## $ PAY_5             : int  -2 0 0 0 0 0 0 0 0 -1 ...
## $ PAY_6             : int  -2 2 0 0 0 0 0 -1 0 -1 ...
## $ BILL_AMT1         : int  3913 2682 29239 46990 8617 64400 367965 11876 11285 0
...
## $ BILL_AMT2         : int  3102 1725 14027 48233 5670 57069 412023 380 14096 0
...
## $ BILL_AMT3         : int  689 2682 13559 49291 35835 57608 445007 601 12108 0
...
## $ BILL_AMT4         : int  0 3272 14331 28314 20940 19394 542653 221 12211 0 ...
## $ BILL_AMT5         : int  0 3455 14948 28959 19146 19619 483003 -159 11793 13007
...
## $ BILL_AMT6         : int  0 3261 15549 29547 19131 20024 473944 567 3719 13912
...
## $ PAY_AMT1          : int  0 0 1518 2000 2000 2500 55000 380 3329 0 ...
## $ PAY_AMT2          : int  689 1000 1500 2019 36681 1815 40000 601 0 0 ...
## $ PAY_AMT3          : int  0 1000 1000 1200 10000 657 38000 0 432 0 ...
## $ PAY_AMT4          : int  0 1000 1000 1100 9000 1000 20239 581 1000 13007 ...
## $ PAY_AMT5          : int  0 0 1000 1069 689 1000 13750 1687 1000 1122 ...
## $ PAY_AMT6          : int  0 2000 5000 1000 679 800 13770 1542 1000 0 ...
## $ default.payment.next.month: int  1 1 0 0 0 0 0 0 0 ...
```

I lowercase the column name, and rename the column names when required, In particular, remarkably this dataset misses a column PAY_1. In the analysis here below we assume that PAY_0 is actually pay_1, to be consider the repayment of the month prior to the month where we calculate the defaulting and removing the duplicate rows.

```
df1 = df

names(df1) = tolower(names(df1))

names(df1)[7] = "pay_1"

# Remove "id" column
df1 <- df1[c(2:25)]

colSums(is.na(df1))
```

```
##          limit_bal          sex
##          0              0
##          education      marriage
##          0              0
##          age            pay_1
##          0              0
##          pay_2          pay_3
##          0              0
##          pay_4          pay_5
##          0              0
##          pay_6          bill_amt1
##          0              0
##          bill_amt2      bill_amt3
##          0              0
##          bill_amt4      bill_amt5
##          0              0
##          bill_amt6      pay_amt1
##          0              0
##          pay_amt2       pay_amt3
##          0              0
##          pay_amt4       pay_amt5
##          0              0
##          pay_amt6 default.payment.next.month
##          0              0
```

```
df1 = df1[!duplicated(df1),]
```

```
cat("Explanatory variables: ", ncol(df1)-1, "\n\n")
```

```
## Explanatory variables: 23
```

```
cat("Number of Observations: ", nrow(df1), "\n\n")
```

```
## Number of Observations: 29965
```

```
df1$default.payment.next.month <- as.factor(df1$default.payment.next.month)
```

```
names(df1)[24] <- "target"
```

```
# create a "target" column for our own convenience
```

```
cat("Target variable: 'default.payment.next.month' -> 'target' \n\n")
```

```
## Target variable: 'default.payment.next.month' -> 'target'
```

Descriptive Analytics

Payment Delays:

Let's start by looking at the past payment delays

```
#names(df1)
```

```
head(df1[c(6:11)],10)
```

```
##      pay_1 pay_2 pay_3 pay_4 pay_5 pay_6
## 1         2     2    -1    -1    -2    -2
## 2        -1     2     0     0     0     2
## 3         0     0     0     0     0     0
## 4         0     0     0     0     0     0
## 5        -1     0    -1     0     0     0
## 6         0     0     0     0     0     0
## 7         0     0     0     0     0     0
## 8         0    -1    -1     0     0    -1
## 9         0     0     2     0     0     0
## 10        -2    -2    -2    -2    -1    -1
```


```
# pay status columns
```

```
library(ggplot2)
par(mfrow=c(3,2))
for(i in 6:11)
{
  hist(df1[,i],main = paste("Histogram of ",names(df1)[i]),labels = FALSE,xlab = names(df1)
[i],col = "green",border = "blue")
}
```


Histogram of pay_1




Histogram of pay_2




Histogram of pay_3




Histogram of pay_4



Histogram of pay_5



Histogram of pay_6



Standing credit

Let's look now at how the debts/credit is accumulating over the months, credit to be repaid is a positive number here.

```
summary(df1[,c(12:17)])
```

```
##      bill_amt1      bill_amt2      bill_amt3      bill_amt4
## Min.      :-165580 Min.      :-69777 Min.      :-157264 Min.      :-170000
## 1st Qu.:   3595 1st Qu.:   3010 1st Qu.:   2711 1st Qu.:   2360
## Median :  22438 Median :  21295 Median :   20135 Median :   19081
## Mean   :   51283 Mean   :   49236 Mean   :   47068 Mean   :   43313
## 3rd Qu.:  67260 3rd Qu.:  64109 3rd Qu.:   60201 3rd Qu.:   54601
## Max.    :  964511 Max.    :  983931 Max.    : 1664089 Max.    :  891586
##      bill_amt5      bill_amt6
## Min.      :-81334 Min.      :-339603
## 1st Qu.:   1787 1st Qu.:    1262
## Median :  18130 Median :   17124
## Mean   :   40358 Mean   :   38917
## 3rd Qu.:  50247 3rd Qu.:   49252
## Max.     :  927171 Max.     :  961664
```

```
head(df1[,c(12:17)],10)
```

```
##      bill_amt1 bill_amt2 bill_amt3 bill_amt4 bill_amt5 bill_amt6
## 1          3913      3102       689         0         0         0
## 2          2682      1725      2682      3272      3455      3261
## 3          29239     14027     13559     14331     14948     15549
## 4          46990     48233     49291     28314     28959     29547
## 5           8617      5670     35835     20940     19146     19131
## 6          64400     57069     57608     19394     19619     20024
## 7          367965    412023    445007    542653    483003    473944
## 8           11876        380        601        221       -159        567
## 9           11285     14096     12108     12211     11793     3719
## 10           0         0         0         0     13007     13912
```

Payments in the previous months

Let's have a quick look at how the payments are performed in the previous month.

```
# pay status columns
```

```
summary(df1[,c(18:23)])
```

```
##      pay_amt1      pay_amt2      pay_amt3      pay_amt4
## Min.   :    0   Min.   :    0   Min.   :    0   Min.   :    0
## 1st Qu.: 1000   1st Qu.:   850   1st Qu.:   390   1st Qu.:   300
## Median : 2102   Median :  2010   Median :  1804   Median :  1500
## Mean   : 5670   Mean   :  5928   Mean   :  5232   Mean   :  4832
## 3rd Qu.: 5008   3rd Qu.:  5000   3rd Qu.:  4512   3rd Qu.:  4016
## Max.   :873552   Max.   :1684259   Max.   :896040   Max.   :621000
##      pay_amt5      pay_amt6
## Min.   :    0   Min.   :    0
## 1st Qu.:   261   1st Qu.:   131
## Median : 1500   Median :  1500
## Mean   : 4805   Mean   :  5222
## 3rd Qu.: 4042   3rd Qu.:  4000
## Max.   :426529   Max.   :528666
```

```
head(df1[,c(18:23)],10)
```

```
##      pay_amt1 pay_amt2 pay_amt3 pay_amt4 pay_amt5 pay_amt6
## 1           0      689         0         0         0         0
## 2           0     1000      1000      1000         0      2000
## 3      1518     1500      1000      1000      1000      5000
## 4      2000     2019      1200      1100      1069      1000
## 5      2000    36681     10000      9000       689       679
## 6      2500     1815       657      1000      1000       800
## 7     55000    40000     38000     20239     13750     13770
## 8        380       601         0       581      1687      1542
## 9      3329         0       432      1000      1000      1000
## 10         0         0         0     13007      1122         0
```

```
summary(df1$limit_bal)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  10000   50000  140000  167442  240000 1000000
```

```
cat("\n Standard deviation:",sd(df1$limit_bal),"\n\n")
```

```
##
## Standard deviation: 129760.1
```

```
# limit balance
```

```
counts <- table(df1$limit_bal)
barplot(counts,log = "y",ylab = "freq",xlab = "Limit Balance")
```



Explore Defaulting


First off, let's start with a zoomed out view on the problem. We want to predict defaulting, Let's answer the following questions:

how many cases do we have on our dataset to work with? What is the breakdown depending on some of the variables available?

```
d = df1

d = table(d$target)

barplot(d,horiz = TRUE,ylab = "target",legend.text = c("Not Default","Default"),col = c("green", "red"))
```



Explore some statistics of defaulting using the categorical variables

Let's have a look at a number of histograms to see how defaulting correlated with the categorical variables available, by converting target, sex, marriage, age to categories

Absolute statistics

```
e = df1

e$target = factor(e$target,levels=c(0,1),labels = c("Not Default","Default"))

e$sex = factor(e$sex,levels=c(1,2),labels = c("Male","Female"))

e$marriage = factor(df1$marriage,levels = c(0,1,2,3),labels = c("na","married","single","other"))

e1 = e %>% group_by(target,sex) %>% summarise(freq = n())

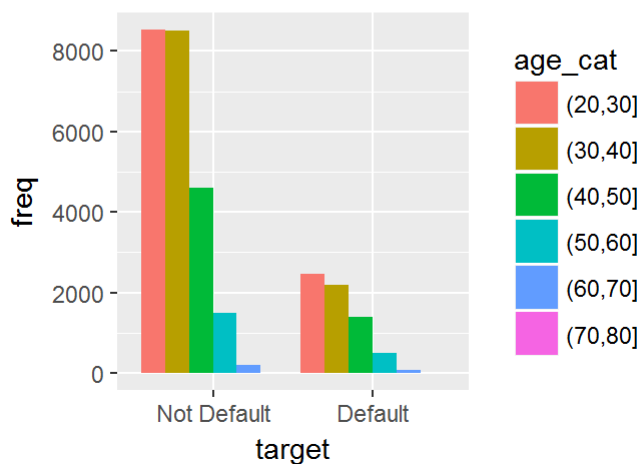


e2 = e %>% group_by(target,marriage) %>% summarise(freq = n())

e$age_cat = cut(e$age,breaks = seq(0,100,10),include.lowest = TRUE)

e3 = e %>% group_by(target,age_cat) %>% summarise(freq = n())

plot1 = ggplot(e1,aes(x=target,y=freq,fill=sex)) + geom_bar(stat = 'identity',position = 'dodge')
plot2 = ggplot(e2,aes(x=target,y=freq,fill=marriage)) + geom_bar(stat = 'identity',position = 'dodge')
plot3 = ggplot(e3,aes(x=target,y=freq,fill=age_cat)) + geom_bar(stat = 'identity',position = 'dodge')

grid.arrange(plot1, plot2,plot3,ncol=2)
```



Statistics relative to the population

```
e = df1

e$target = factor(e$target,levels=c(0,1),labels = c("Not Default","Default"))

e$sex = factor(e$sex,levels=c(1,2),labels = c("Male","Female"))

e$marriage = factor(df1$marriage,levels = c(0,1,2,3),labels = c("na","married","single","other"))

e1 = e %>% group_by(target,sex) %>% summarise(freq = n())
e11 = e %>% group_by(sex) %>% summarise(freq1 = n())
e1$rel_freq = e1$freq/e11$freq1

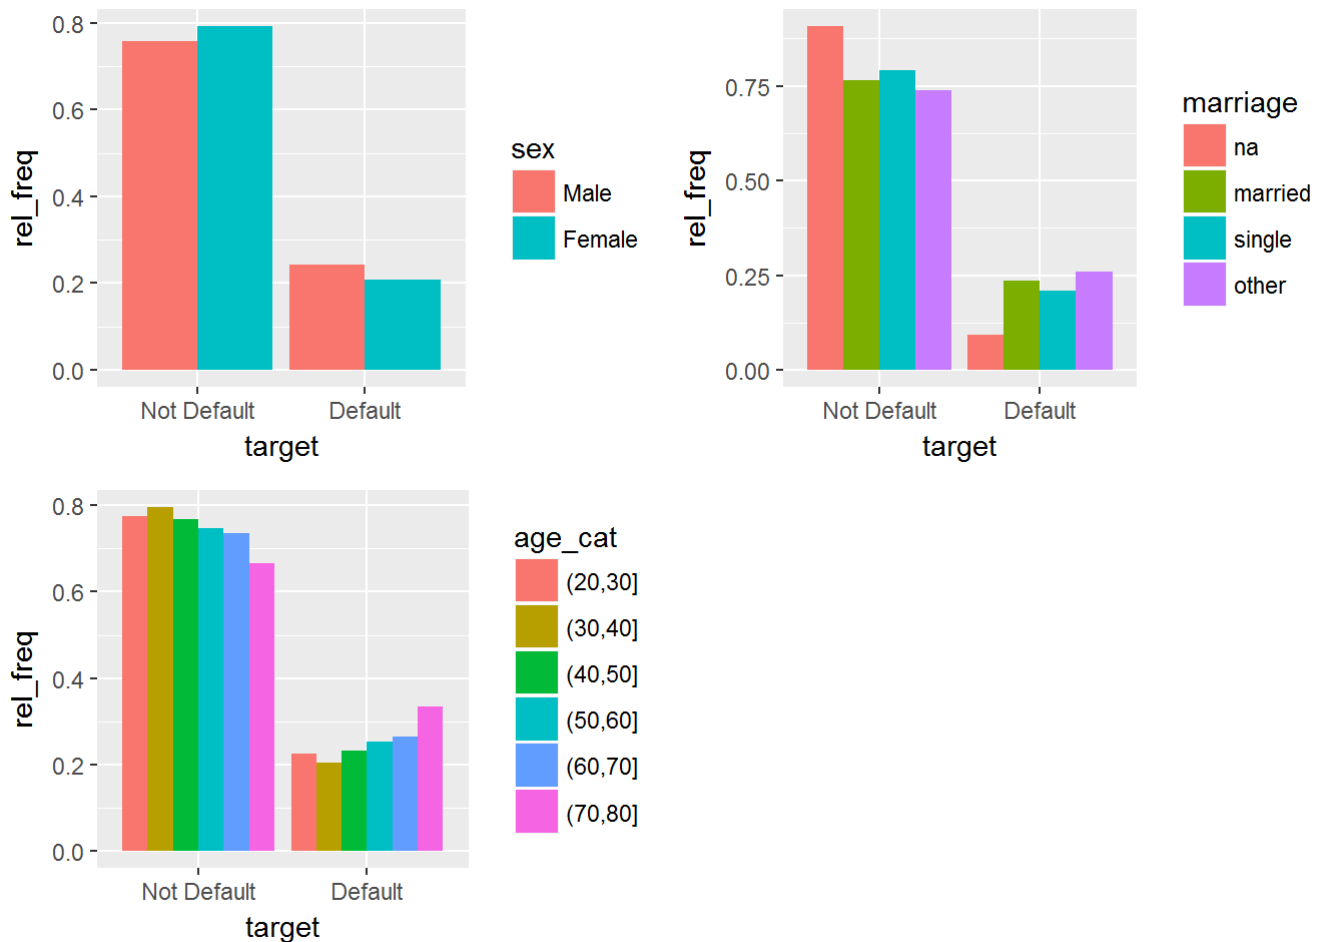
e2 = e %>% group_by(target,marriage) %>% summarise(freq = n())
e11 = e %>% group_by(marriage) %>% summarise(freq1 = n())
e2$rel_freq = e2$freq/e11$freq1

e$age_cat = cut(e$age,breaks = seq(0,100,10),include.lowest = TRUE)

e3 = e %>% group_by(target,age_cat) %>% summarise(freq = n())
e11 = e %>% group_by(age_cat) %>% summarise(freq1 = n())
e3$rel_freq = e3$freq/e11$freq1

plot1 = ggplot(e1,aes(x=target,y=rel_freq,fill=sex)) + geom_bar(stat = 'identity',position = 'dodge')
plot2 = ggplot(e2,aes(x=target,y=rel_freq,fill=marriage)) + geom_bar(stat = 'identity',position = 'dodge')
plot3 = ggplot(e3,aes(x=target,y=rel_freq,fill=age_cat)) + geom_bar(stat = 'identity',position = 'dodge')

grid.arrange(plot1, plot2,plot3,ncol=2)
```



Feature engineering

Splitting the dataset into the Training set and Test set

```
d = df1

#d[,c(1,5,12:23)] = scale(d[,c(1,5,12:23)])

library(caTools)
```

```
## Warning: package 'caTools' was built under R version 3.4.2
```

```
set.seed(123)
split = sample.split(d, SplitRatio = 0.7)

train = subset(d, split==T)
test = subset(d, split==F)

train[,c(1,5,12:23)] = scale(train[,c(1,5,12:23)])
test[,c(1,5,12:23)] = scale(test[,c(1,5,12:23)])
```

Models

Support Vector Machine (SVM)

```
library(e1071)
```

```
## Warning: package 'e1071' was built under R version 3.4.2
```

```
model1 = svm(formula = target ~ ., data = train, type = 'C-classification', kernel = 'linear'
)

prob_pred = predict(model1,newdata = test[, -24])

cm = table(prob_pred, test[, 24])

# cm = table(prob_pred, test[, 10])

accuracy = sum(diag(cm)) / sum(cm)

cat("SVM model's accuracy: ", accuracy)
```

```
## SVM model's accuracy: 0.8068682
```

Logistic Regression

```
model2 = glm(formula = target ~ ., data = train, family = binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
#model2 = glm(formula = target ~ limit_bal+education+marriage+age+pay_1+pay_2+pay_3+bill_amt1
+pay_amt1+pay_amt2, data = train, family = binomial)

prob_pred = predict(model2, newdata = test[, -24], type = 'response')

y_pred = ifelse(prob_pred > 0.5, 1, 0)

cm = table(y_pred, test[, 24])

# cm = table(prob_pred, test[, 10])

accuracy = sum(diag(cm)) / sum(cm)

cat("Logistic Regression model's accuracy: ", accuracy)
```

```
## Logistic Regression model's accuracy: 0.8129756
```

```
#str(d)
```

Naive Bayes


```
library(e1071)

model3 = naiveBayes(formula = target ~ ., data = train)

prob_pred = predict(model3,newdata = test[, -24])

cm = table(prob_pred, test[, 24])

# cm = table(prob_pred, test[, 10])

accuracy = sum(diag(cm)) / sum(cm)

cat("Naive Bayes model's accuracy: ", accuracy)
```

```
## Naive Bayes model's accuracy: 0.7300761
```

Decision Tree

```
library(rpart)

model4 = rpart(formula = target ~ ., data = train)

prob_pred = predict(model4, newdata = test[, -24], type = 'class')

cm = table(prob_pred, test[, 24])

# cm = table(prob_pred, test[, 10])

accuracy = sum(diag(cm)) / sum(cm)

cat("Decision Tree Classification model's accuracy: ", accuracy)
```

```
## Decision Tree Classification model's accuracy: 0.8201842
```